

## Assignment 1: Probability Distributions in R

Many statistical tools and techniques used in data analysis are based on probability. Probability measures how likely it is for an event to occur on a scale from 0 (the event never occurs) to 1 (the event always occurs.). When working with data, variables in the columns of the data set can be thought of as random variables: variables that vary due to chance. A [probability distribution](#) describes how a random variable is distributed; it tells us which values a random variable is most likely to take on and which values are less likely.

In statistics, there are a range of precisely defined probability distributions that have different shapes and can be used to model different types of random events. In this lesson we'll discuss some common probability distributions and how to work with them in R.

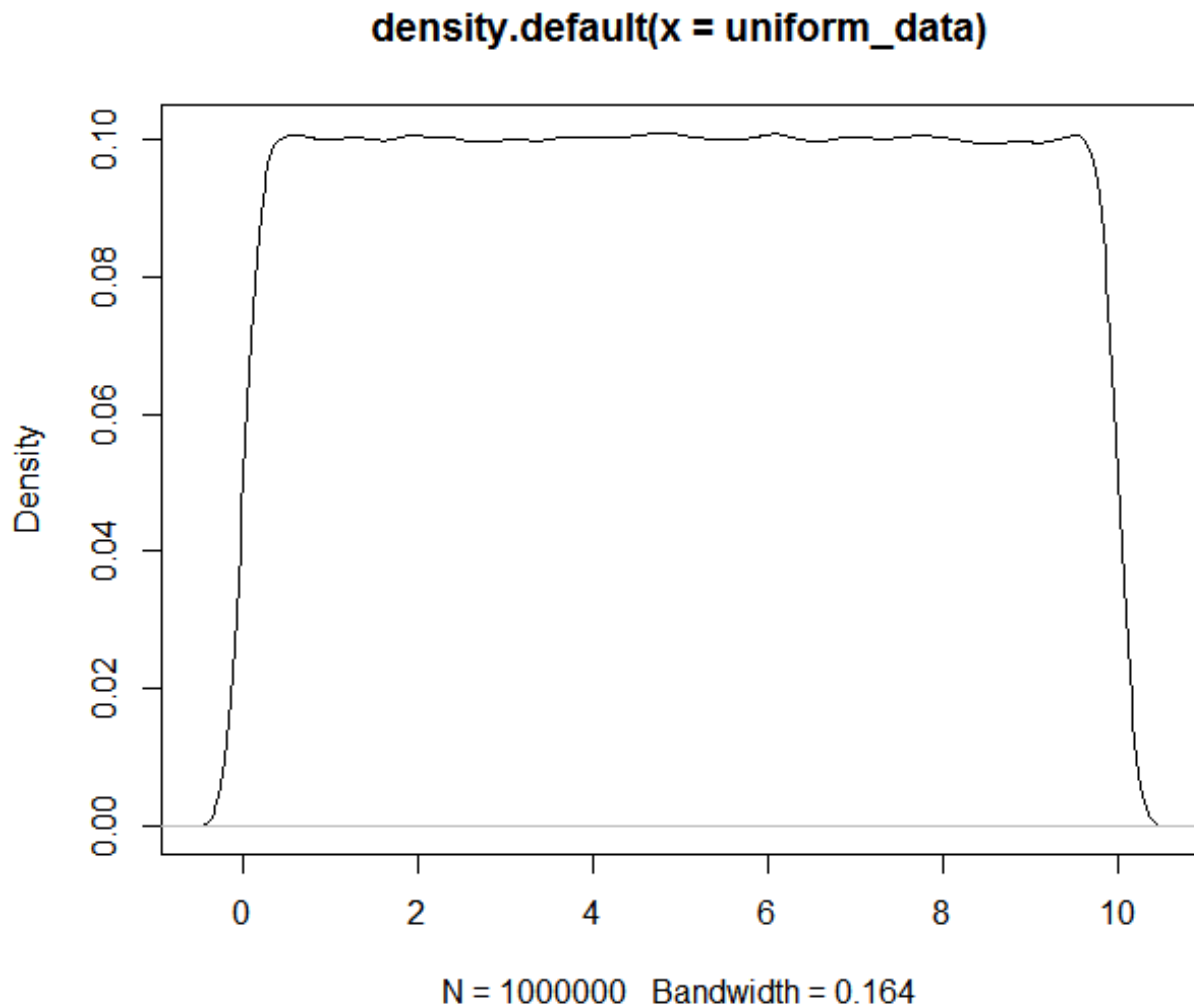
### Distribution Function Basics and The Uniform Distribution

The [uniform distribution](#) is a probability distribution where each value within a certain range is equally likely to occur and values outside of the range never occur. If we make a density plot of a uniform distribution, it appears flat because no value is any more likely (and hence has any more density) than another:

In [1]:

```
uniform_data <- runif(1000000,0,10)    # generate data from a uniform distribution
```

```
plot(density(uniform_data ))           # create a density plot
```



**\*Note:** the plot above is an approximation of the underlying distribution, since it is based on a sample of observations.

In the code above, we generated 1 million data points from a uniform distribution spanning the range 0 to 10. In the density plot we see that the density of our uniform data is essentially level meaning any given value has the same probability of occurring. The area under a probability density curve is always equal to 1.

Every probability distribution in R comes with 4 function prefixes: 'r', 'p', 'q' and 'd'. When working probability distributions in R, your function calls will start with one of these four prefixes followed by a name for the type distribution you're working with. For instance, the R name for the uniform distribution is 'unif' so the uniform distribution functions available to us are `runif()`, `punif()`, `qunif()` and `dunif()`. Let's go over each of these prefixes in detail.

The 'r' prefix is used to generate random numbers from the specified distribution. As you may have noticed in previous lessons, we've used the `runif()` function many times generate random numbers. The arguments you supply when generating random numbers depend on the type of distribution you're

using. In the case of the uniform, we have to specify how many random numbers we want to generate and then the range of the distribution:

In [2]:

```
more_random_data <- runif(n=10000,      # Generate 10,000 numbers
                          min=0,        # In the range 0
                          max=10)       # To 10
```

The prefix 'p' is used to determine the probability that an observation drawn from a distribution falls below a specified value (known as the [cumulative distribution function](#)). In essence, 'p' gives you the area under the distribution's density curve to the left of a certain value on the x axis. For example, in the uniform distribution above, there is a 25% chance that an observation will be in the range 0 to 2.5 and a 75% chance it will fall in the range 2.5 to 10. We can confirm this with punif():

In [3]:

```
punif(q = 2.5,          # Cutoff value (quantile) to check
      min=0,
      max=10)
```

Out[3]:

0.25

The prefix 'q' is the inverse of the prefix 'p': it returns the cutoff value (quantile) associated with a given probability. For instance, if we want to know the cutoff value for which we have a 40% chance of drawing an observation below that value, we can use q:

In [4]:

```
qunif(p = 0.4,          # Probability cutoff
      min = 0,
      max = 10)
```

Out[4]:

4

Finally the 'd' prefix gives you the density (height of the density curve) at a given point. Our plot of randomly generated uniform data was flat at 0.1 so the density of the distribution should be constant at 0.1 throughout the range 0 to 10:

In [5]:

```
dunif(x = 0,            # x value at which to check density
      min = 0,
      max = 10)
```

```
dunif(5, 0, 10)         # Check some more densities
```

```
dunif(10, 0, 10)
```

```
dunif(11, 0, 10)        # a value outside the distribution's range has 0 probability
```

Out[5]:

0.1

## Generating Random Numbers and Setting The Seed

Any time you need to generate random real numbers in a range with equal probability you can use the `runif()` function. The uniform distribution is continuous, meaning the values it generates can take on any real numbered value within the specified range, so the numbers it generates are decimals. If you need to generate integers instead of decimals, you could use `floor()` to round the decimals down:

In [6]:

```
random_ints <- floor(runif(100000,1,11))    # Generate random data and round down
```

```
table(random_ints)                        # Check counts
```

Out[6]:

```
random_ints
  1     2     3     4     5     6     7     8     9    10
10028  9976  9916  9914  9919 10076 10055 10076 10050  9990
```

Using `runif()` is not necessary to generate integers however: you can use the `sample()` function to draw a specified number of samples from a vector instead:

In [7]:

```
random_ints2 <- sample(x = 1:10,           # Vector to sample from
                      size = 100000,       # Number of samples
                      replace = TRUE)      # Sample with replacement?*
```

```
table(random_ints2)                      # Check counts
```

Out[7]:

```
random_ints2
  1     2     3     4     5     6     7     8     9    10
10115  9952  9896  9905 10190  9924 10103  9929 10014  9972
```

**\*Note:** Sampling with replacement means you replace or "put back" a value after drawing it instead of removing it from the set you are sampling.

When generating random data and working with functions that involve randomization, your results can differ from one run to the next. Having results vary each time you run a function is often not desirable. For example, if you want a colleague to be able to reproduce your results exactly, you can run into problems when you use randomization. You can ensure that your results are the same each time you use a function that involves randomness by setting the random number generator's seed value to initialize it prior to running the function. Set the seed with `set.seed()`:

In [8]:

```
set.seed(12)    # Set seed to arbitrary value
```

```
runif(10,1,10)  # Generate some random numbers
```

```
set.seed(12)    # Reset seed to the same value
```

```
runif(10,1,10)  # Generate more random numbers
```

Out[8]:

```
1. 1.62424824223854
```

2. 8.35997678851709

3. 9.4835955912713

4. 3.4244368870277

5. 2.52413310925476

6. 1.30506060225889

7. 2.60906503768638

8. 6.77498828992248

9. 1.2058996912092

10. 1.07492343941703

11. 1.62424824223854

12. 8.35997678851709

13. 9.4835955912713

14. 3.4244368870277

15. 2.52413310925476

16. 1.30506060225889

17. 2.60906503768638

18. 6.77498828992248

19. 1.2058996912092

20. 1.07492343941703

Notice that we generated the exact same numbers with both calls to `runif()` because we set the same seed before each call. If we had not set the seed, we would have gotten different numbers. This reproducibility illustrates the fact that these random numbers aren't truly random, but rather "[pseudorandom](#)".

## The Normal Distribution

The normal or [Gaussian distribution](#) is a continuous probability distribution characterized by a symmetric bell-shaped curve. A normal distribution is defined by its center (mean) and spread (standard deviation.). The bulk of the observations generated from a normal distribution lie near the mean, which lies at the exact center of the distribution: as a rule of thumb, about 68% of the data lies within 1 standard deviation of the mean, 95% lies within 2 standard deviations and 99.7% lies within 3 standard deviations.

The normal distribution is perhaps the most important distribution in all of statistics. It turns out that many real world phenomena, like IQ test scores and human heights, roughly follow a normal distribution, so it is often used to model random variables. Many common statistical tests assume distributions are normal. The R nickname for the normal distribution is 'norm' so the functions for working with it are: `rnorm()`, `pnorm()`, `qnorm()` and `dnorm()`:

In [9]:

```
normally_distributed <- rnorm(1000000,      # Generate normally distributed data
                              mean = 0,      # With mean 0
                              sd = 1)       # And standard deviation 1

prob_under_minus1 <- pnorm(q=-1,           # Get prob of observing a value less than -1
                           mean=0,
                           sd=1)

prob_over_1 <- 1-pnorm(q=1,                # Get prob of observing a value over 1
                      mean=0,
                      sd=1)

between_prob <- 1-(prob_under_minus1+prob_over_1) # Prob between -1 and 1
```

```
prob_under_minus1
prob_over_1
between_prob
```

Out[9]:

**0.158655253931457**

Out[9]:

**0.682689492137086**

The output shows that roughly 16% of the data generated by a normal distribution with mean 0 and standard deviation 1 is below -1, 16% is above 1 and 68% lies between -1 and 1, which agrees with the 68, 95, 99.7 rule. Let's plot the randomly generated normal data and inspect the three areas we calculated under the density curve:

In [10]:

```
library(ggplot2)
```


```

In [11]:
# Plot the density curve with the cutoff areas
norm_frame = with(density(normally_distributed), # Create data frame density
values

                    data.frame(x,y))

myplot <- ggplot(data = norm_frame, aes(x = x, y = y)) + # Create the plot
  geom_line() +
  geom_ribbon(data=subset(norm_frame,x < -1),
            aes(ymax=y),
            ymin=0,
            fill="red",
            alpha=0.4) +
  geom_ribbon(data=subset(norm_frame,x > 1),
            aes(ymax=y),
            ymin=0,
            fill="red",
            alpha=0.4) +
  geom_ribbon(data=subset(norm_frame,x > -1 & x < 1),
            aes(ymax=y),
            ymin=0,
            fill="skyblue",
            alpha=0.4) +
  geom_text(x=-1.6,y=0.03,label=round(prob_under_minus1,4),size=4) +
  geom_text(x=1.6,y=0.03,label=round(prob_under_minus1,4),size=4) +
  geom_text(x=0,y=0.1,label=round(1-(prob_under_minus1*2),4),size=5) +
  xlim(-4,4)

```



The plot above shows the bell shape of the normal distribution, the area below and above one standard deviation and the area within 1 standard deviation of the mean.

Finding quantiles of the normal distribution is a common task when performing statistical tests. You can check quantiles with `qnorm()`:

In [12]:

```
qnorm(p = 0.025)      # Find the quantile at the 2.5% cutoff
qnorm(p = 0.975)      # Find the quantile at the 97.5% cutoff
```

Out[12]:

**-1.95996398454005**

Out[12]:

**1.95996398454005**

The quantile output above confirms that roughly 5% of the data lies more than 2 standard deviations from the mean.



**\*Note: a mean of 0 and standard deviation of 1 are default values for the normal distribution functions.**

## The Binomial Distribution

The [binomial distribution](#) is a discrete probability distribution that models the outcomes of a given number of random trials of some experiment or event. The binomial is defined by two parameters: the probability of success in any given trial and the number of trials. The binomial distribution tells you how likely it is to achieve a given number of successes in  $n$  trials of the experiment. For example, we could model flipping a fair coin 10 times with a binomial distribution where the number of trials is set to 10 and the probability of success is set to 0.5. In this case the distribution would tell us how likely it is to get zero heads, 1 head, 2 heads and so on.

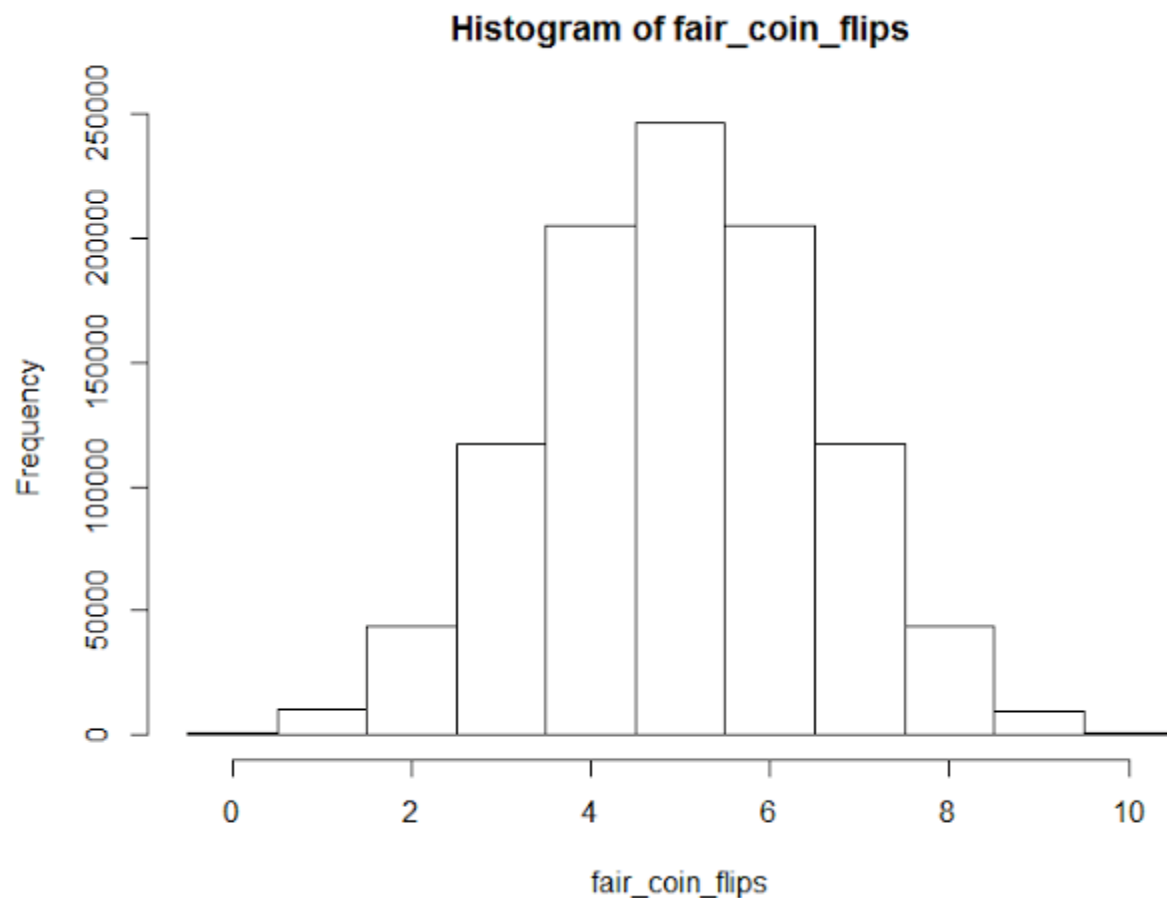
The R nickname for the binomial distribution is 'binom', so the functions for working with it are `rbinom()`, `dbinom()`, `qbinom()` and `pnbinom()`. Let's generate and investigate some binomial data:

In [13]:

```
fair_coin_flips <- rbinom(1000000,      # Generate data from the binomial distribution
                        size = 10,     # With 10 trials
                        prob = 0.5)    # And success probability 0.5

table(fair_coin_flips)                # Check the counts

hist(fair_coin_flips, breaks=seq(-0.5,10.5,1)) # Plot the results as a histogram
```



Out[13]:

```
fair_coin_flips
  0      1      2      3      4      5      6      7      8      9     10
973   9854  43735 117091 204867 246727 204975 117295  43722  9786   975
```

**Note that since the binomial distribution is discrete, it only takes on integer values so we can summarize binomial data with a table and its distribution with a histogram. The histogram shows us that a binomial distribution with a 50% probability of success is symmetric, with the most likely outcomes lying at the center. This is reminiscent of the normal distribution, but if we alter the success probability, the distribution won't be symmetric:**

In [14]:

```
set.seed(9)
biased_coin_flips <- rbinom(1000000, # Generate data from the binomial dist
                             size = 10, # With 10 trials
                             prob = 0.8) # And success probability 0.8

table(biased_coin_flips) # Check the counts
```

```
# hist(biased_coin_flips, breaks=seq(-0.5,10.5,1))    # Plot the results as a histogram
```

Out[14]:

biased\_coin\_flips

| 1 | 2  | 3   | 4    | 5     | 6     | 7      | 8      | 9      | 10     |
|---|----|-----|------|-------|-------|--------|--------|--------|--------|
| 1 | 60 | 756 | 5395 | 26336 | 88095 | 201120 | 302086 | 269178 | 106973 |

**The pbinom() function lets us check the probability of achieving a number of successes within a certain range:**

In [15]:

```
pbinom(q=5,                                # Check the chances of achieving 5 successes or less
       size=10,
       prob=0.8)
```

```
pbinom(q=8,                                # Check the chances of achieving 9 successes or more
```

```

        size=10,
        prob=0.8,
        lower.tail = FALSE)      # Check the upper tail instead of the lower tail

```

Out[15]:

**0.0327934976**

Out[15]:

**0.375809638400001**

**dbinom()** checks the density of the binomial at a given *x* value. Since the binomial is discrete, this is the probability of getting a specific number of successes:

In [16]:

```

# Check the probability of 5 successes in 10 flips of a fair coin
dbinom(x=5,
        size=10,
        prob=0.5)

```

```

# Check the probability of 8 successes in 10 flips of a biased coin
dbinom(x=8,
        size=10,
        prob=0.8)

```

Out[16]:

**0.24609375**

Out[16]:

**0.301989888**

## The Geometric and Exponential Distributions

The [geometric](#) and [exponential](#) distributions model the time it takes for an event to occur. The geometric distribution is discrete and models the number of trials it takes to achieve a success in repeated experiments with a given probability of success. The exponential distribution is a continuous analog of the geometric distribution and models the amount of time you have to wait before an event occurs given a certain occurrence rate.

The R name for the geometric distribution is 'geom'. Let's use the geom functions to model the number of trials it takes to get a success (heads) when flipping a fair coin:

In [17]:

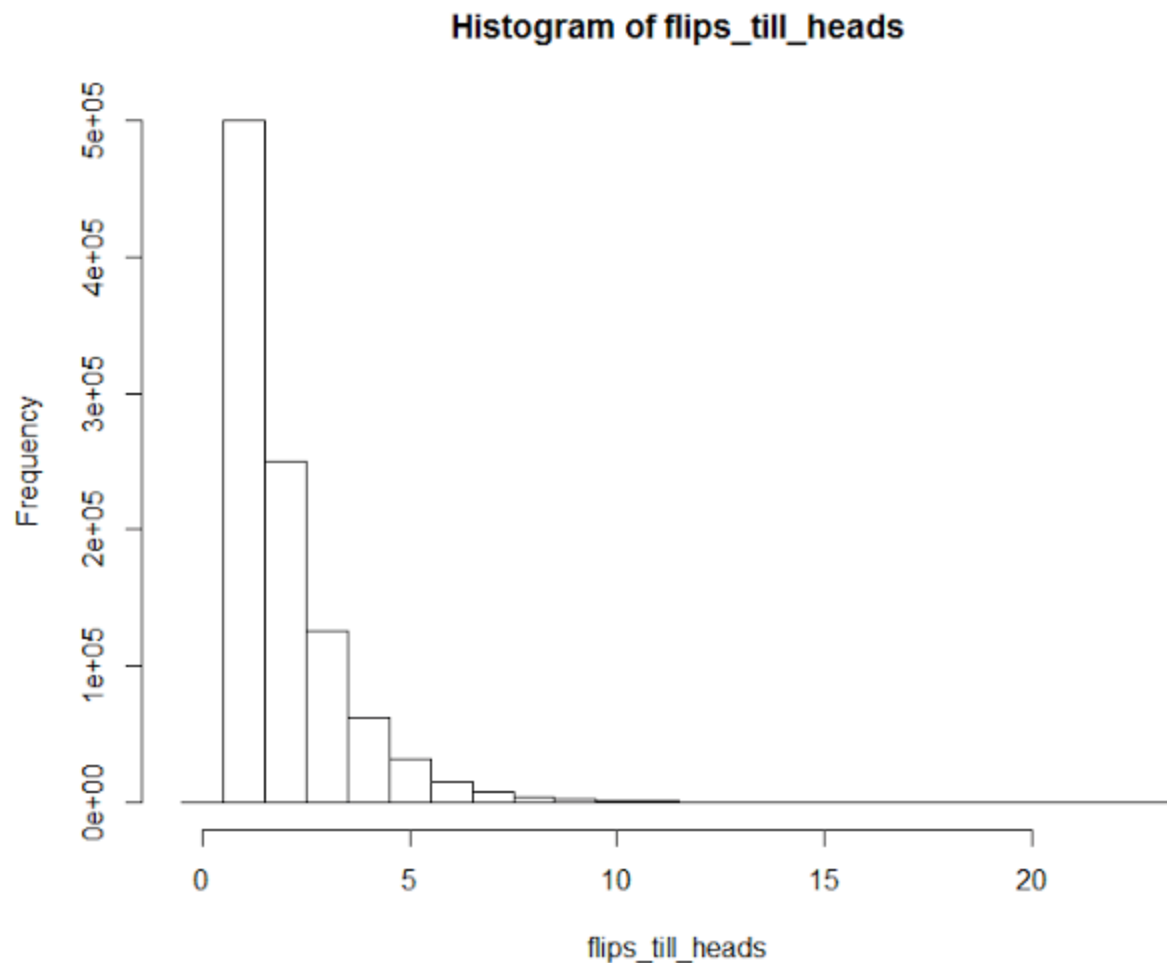
```

set.seed(12)
flips_till_heads <- rgeom(n=1000000,      # Generate geometric data*
                          prob=0.5) + 1

table(flips_till_heads)      # Check counts of results

hist(flips_till_heads, breaks=seq(-0.5,max(flips_till_heads)+0.5))  # Plot the results

```



Out[17]:

flips\_till\_heads

|        |        |        |       |       |       |      |      |      |     |     |
|--------|--------|--------|-------|-------|-------|------|------|------|-----|-----|
| 1      | 2      | 3      | 4     | 5     | 6     | 7    | 8    | 9    | 10  | 11  |
| 500135 | 250046 | 125003 | 62229 | 31088 | 15649 | 7976 | 3869 | 2044 | 976 | 481 |
| 12     | 13     | 14     | 15    | 16    | 17    | 18   | 19   | 20   | 22  | 23  |
| 256    | 123    | 66     | 28    | 15    | 7     | 4    | 1    | 2    | 1   | 1   |

**\*Note:** in R the geom functions actually models the number of failures until the first success, so we need add 1 to get the number of trials until the first success.

The distribution looks similar to what we'd expect: it is very likely to get a heads in 1 or 2 flips, while it is very unlikely for it to take more than 5 flips to get a heads. In the 1 million trials we generated, the longest it took to get a heads was 23 flips.

Let's use pgeom() to check the probability of needing 6 flips or more to get a success:

In [18]:

```
pgeom(q=4,
      prob=0.5,
      lower.tail = FALSE)    # Check the upper tail
```

Out[18]:

0.03125

**You can use `dgeom()` to check the probability of seeing a specific number of failures before a success:**

In [19]:

```
dgeom(x=1,          # Check the chances of getting 1 failure and then a success
      prob=0.5)
```

Out[19]:

0.25

**The R name for the exponential distribution is 'exp'. We've already used the exponential function a few times in past lessons to generate skewed data with `rexp()`. Let's generate some exponential data, calculate the probability that we have to wait longer than 1 time unit for a success and plot it:**

In [20]:

```
exponential_data <- rexp(n=1000000,
                        rate=1)          # Success/arrival rate*

# Get the probability of waiting more than 1 time unit before a success
prob_longer_1 <- pexp(q=1,
                    rate=1,
                    lower.tail=FALSE)
```

prob\_longer\_1

Out[20]:

0.367879441171442


**\*Note: The average arrival time for the exponential distribution is equal to  $1/\text{rate}$ .**

In [21]:

```
# Plot the density curve of the exponential
exp_frame = with(density(exponential_data), # Create data frame of x,y density values
                data.frame(x,y))

myplot <- ggplot(data = exp_frame, aes(x = x, y = y)) + # Create the plot
  geom_line() +
  geom_ribbon(data=subset(exp_frame,x > 1),
            aes(ymax=y),
            ymin=0,
            fill="red",
            alpha=0.4)+
  geom_ribbon(data=subset(exp_frame,x < 1),
            aes(ymax=y),
            ymin=0,
            fill="skyblue",
            alpha=0.4) +
```

```
geom_text(x=2,y=0.06,label=round(prob_longer_1,4), size=5) +
geom_text(x=0.5,y=0.125,label=round(1-prob_longer_1,4), size=5) +
xlim(-0.5,7)
```



Similar to the geometric distribution, the exponential starts high and has a long tail that trails off to the right that contains rare cases where you have to wait much longer than average for an arrival.

## The Poisson Distribution

The [Poisson distribution](#) models the probability of seeing a certain number of successes within a time interval, where the time it takes for the next success is modeled by an exponential distribution. The Poisson distribution can be used to model traffic, such as the number of arrivals a hospital can expect in a hour's time or the number of emails you'd expect to receive in a week.


The R shorthand for the Poisson distribution is 'pois'. Let's generate and plot some data from a Poisson distribution with an arrival rate of 1:

In [22]:

```
set.seed(12)
arrival_rate_1 <- rpois(n=1000000,
                       lambda=1)      # Arrival rate

table(arrival_rate_1)

hist(arrival_rate_1,breaks=seq(-0.5,max(arrival_rate_1)+0.5,1))
```



Out[22]:

```
arrival_rate_1
  0      1      2      3      4      5      6      7      8      9
367994 367884 183697  61524  15193   3096   527   70   13    2
```



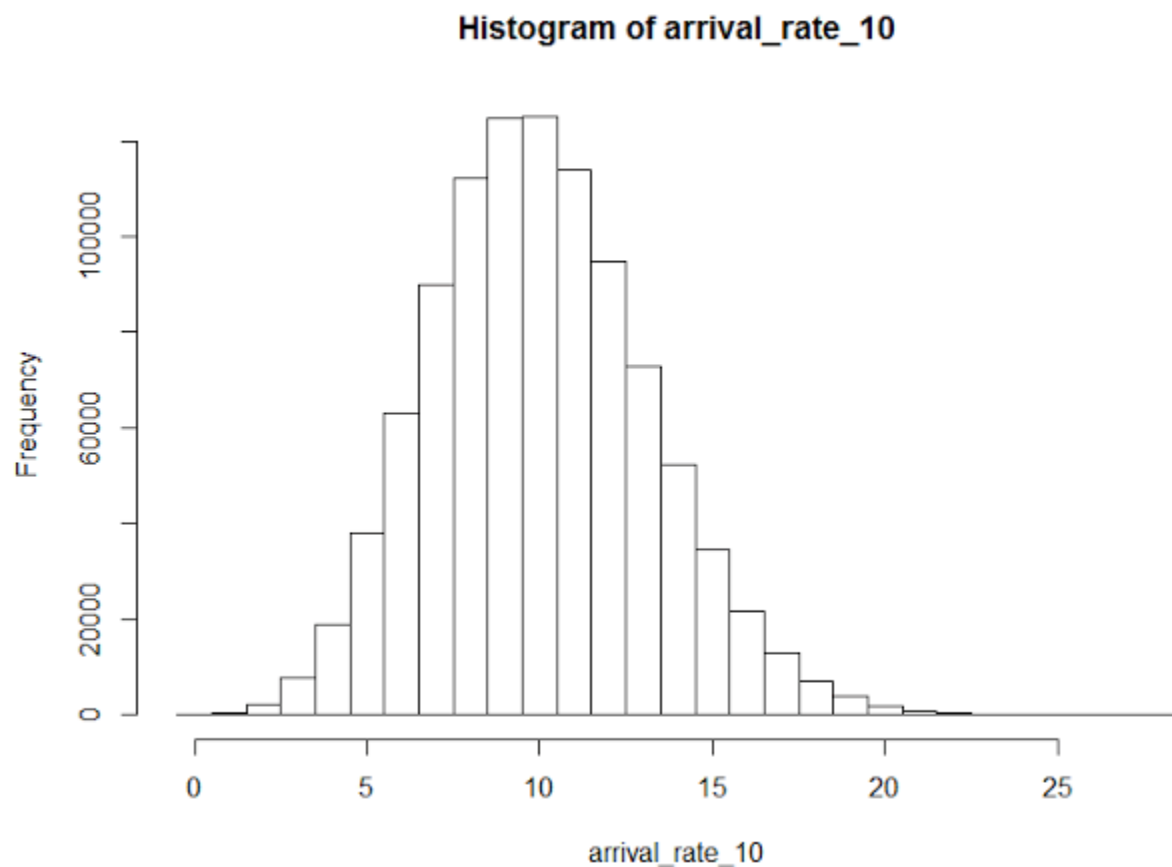
The histogram shows that when arrivals are relatively infrequent, it is rare to see more than a couple of arrivals in each time period. When the arrival rate is high, it becomes increasingly rare to see a low number of arrivals and the distribution starts to look more symmetric:

In [23]:

```
set.seed(12)
arrival_rate_10 <- rpois(n=1000000,
                        lambda=10)           # Arrival rate of 10 per time unit

table(arrival_rate_10)

hist(arrival_rate_10,breaks=seq(-0.5,max(arrival_rate_10)+0.5,1))
```



Out[23]:

```
arrival_rate_10
  0      1      2      3      4      5      6      7      8      9     10
49    473   2222   7576  18811  38135  62970  89986 112335 124749 125202
11    12    13    14    15    16    17    18    19    20    21
114003 94613 72881 52293 34696 21709 12811 7147 3818 1952 874
```

|     |     |    |    |    |    |    |
|-----|-----|----|----|----|----|----|
| 22  | 23  | 24 | 25 | 26 | 27 | 28 |
| 400 | 175 | 65 | 30 | 12 | 10 | 3  |

We can use `ppois()` to check the probability of achieving more or less than a certain number of successes and `dpois()` to check the probability of obtaining a specific number of successes:

In [24]:

```
ppois(q=5,          # check the probability of getting 5 successes or less
      lambda=10)    # With an arrival rate of 10
```

```
dpois(x=10,        # Check the probability of getting exactly 10 successes
      lambda=10)
```

Out[24]:

**0.0670859628790318**

Out[24]:

**0.125110035721133**

## Wrap Up

R contains built in functions that make it easy to work with a wide range of probability distributions, including [many that we did not discuss](#) in this lesson. Probability distribution functions are useful for generating random data, modeling random events and aiding with statistical tests and analysis.

In the next few lessons we'll learn how to carry out common statistical tests in R.